



A neural meta model for predicting winter wheat crop yield

Yogesh Bansal¹ · David Lillis¹ · M.-Tahar Kechadi^{1,2}

Received: 12 June 2023 / Revised: 21 August 2023 / Accepted: 10 October 2023

© The Author(s), under exclusive licence to Springer Science+Business Media LLC, part of Springer Nature 2024

Abstract

This study presents the development and evaluation of machine learning models to predict winter wheat crop yield using heterogeneous soil and weather data sets. A concept of an error stabilisation stopping mechanism is introduced in an LSTM model specifically designed for heterogeneous datasets. The comparative analysis of this model against an LSTM model highlighted its superior predictive performance. Furthermore, weighted regression models were developed to capture environmental factors using agroclimatic indices. Finally, a neural meta model was built by combining the predictions of several individual models. The experimental results indicated that a neural meta model with an MAE of 0.82 and RMSE of 0.983 tons/hectare demonstrated a notable performance, highlighting the importance of incorporating weighted regression models based on agroclimatic indices. This study shows the potential for improved yield prediction through the proposed model and the subsequent development of a meta model.

Keywords Neural meta model · Soil and weather dataset · Agro-climatic indices · Feature engineering · LSTM · Weighted regression ML models

1 Introduction

Crop production plays a key role in ensuring food security and supporting economies around the world. An extensive number of studies have been conducted in this field using a variety of machine learning (ML) models, covering deep learning and ensemble models on diverse datasets to improve crop yield predictions. However, a recurring challenge in this

Editor: Vu Nguyen, Dani Yogatama.

✉ Yogesh Bansal
yogesh.bansal@ucdconnect.ie

David Lillis
david.lillis@ucd.ie

M.-Tahar Kechadi
tahar.kechadi@ucd.ie

¹ School of Computer Science, University College Dublin, Dublin 4, Ireland

² Insight Centre for Data Analytics, Dublin 4, Ireland

field is the effective management of heterogeneous data sets, which can lead to less optimal results. Motivated by this challenge, our aim is to improve a previously built long short-term memory (LSTM) model specifically designed to handle such datasets. Previously, we established an LSTM (Bansal et al., 2023) model designed specifically for heterogeneous data sets. In this paper, we advance our model by introducing error stabilisation stopping within LSTM with the aim of improving the accuracy of yield predictions.

The presence of agroclimatic indices is important to capture the environmental conditions on a zone level. To do this, we build weighted regression ML models (decision tree, random forest, and gradient boosting). Lastly, we build several neural meta models that combines the strengths of individual models. The aim of this study is to provide an improved machine learning model for the prediction of winter wheat crop yield.

The paper is structured as follows. The literature is presented in Sect. 2, followed by a description of the data in Sect. 3. Pre-processing information is provided in Sect. 4, followed by “experiment setup” and “results” in Sects. 6 and 7, respectively. Section 9 concludes the study.

2 Literature

The incorporation of machine learning (ML), including deep learning techniques, into agricultural crop yield predictions has garnered considerable attention in recent years. These data-driven methodologies play an increasingly vital role in shaping the direction of research in this field. The increasing interest is reflected in the volume of studies and surveys that have rigorously explored, evaluated, and compared various predictive models.

Elavarasan et al. (2018) began by surveying various machine learning models and comparing their performance using error measures. Expanding on this, Klompenburg et al. (2020) provided a systematic review of the applications of machine learning in the prediction of crop yields. Building on these initial works, Oikonomidis et al. (2022a) focused on deep learning methods, demonstrating their potential to improve prediction accuracy. Bali & Singla (2022) explored emerging trends, underscoring the influence of various factors on crop yield. More recently, Toomula & Pelluri (2021) and Muruganantham et al. (1990) offered updated views on deep learning-based models, the latter identifying LSTM and CNN as the most effective techniques, thus signifying continuous innovation in this rapidly evolving field.

Hybrid models, which combine various ML models, have demonstrated superior performance in crop yield predictions compared to individual ML models. The study by Sun et al. (2019) provides a notable example, where the authors proposed deep convolutional neural networks (CNN) for the prediction of soybean yield at the end of the season and at the county level. Their findings indicated that the proposed CNN-LSTM model outperformed either the CNN or LSTM model alone, suggesting that such hybrid model architectures have considerable potential to improve yield predictions for various other crops, including corn, wheat, and potatoes. This sentiment echoes the research by Wang et al. (2020), who developed an LSTM-CNN model to estimate winter wheat yield at the county level in China. The LSTM-CNN model was based on weather and remote sensing data, and the results indicated that it significantly improved the model’s yield prediction capabilities.

Agroclimatic indices also play a crucial role in crop yield modelling. As Mathieu and Aires (2018) highlighted, these indices, including factors such as temperature and precipitation, can significantly improve crop yield modelling compared to using direct weather

variables alone. This understanding of the agroclimatic conditions of a region can enable farmers to better manage their crops, improve yields, and minimise the risks associated with climate variability. Chergui et al. (2022) provide a comprehensive review of machine learning models applied to crop management, shedding light on various ML techniques and their effectiveness in different contexts. Wang et al. research (2020) further emphasises the potential of ML, where they successfully used neural models to estimate winter wheat production based on soil and weather data, and the neural models outperform other models.

Sharma et al. (2011) used an LSTM model to predict agricultural yields through satellite data in India, outperforming other ML techniques. A similar superiority of LSTM was observed in studies by Iniyan et al. (2023) and Bali and Singla (2021). Furthermore, combining LSTM with random forest showed an improved forecast accuracy, as shown by Shen et al. (2022). Bansal et al. (2022, 2023) highlighted the importance of weather data and introduced a deep learning model to refine the predictions of the winter wheat crop yield.

Shahhosseini et al.'s research (2020, 2021) centred on ensemble machine learning models, with CNN-DNN ensembles presenting accurate results. Dang et al. (2021) compared various models, finding Deep Neural Networks to be the most effective. Pravallika et al. (2021) and Srivastava et al. emphasised the significant impact of environmental variables on yield predictions (Srivastava et al., 2022). Zhang et al. advocate for combining deep learning and ensemble approaches (Zhang et al., 2022a, 2022b), and the integration of various data sources, such as satellite imagery, has been explored in numerous studies (Oikonomidis et al., 2022b; Sarith Divakar et al., 2022). Hybrid deep learning frameworks like that of Khaki et al. have shown promising outcomes (Khaki et al., 1750).

3 Data description

The data sources used in this study consist of soil and weather data sets that cover a large number of farms for several years. A farm is divided into several fields, and a field is further divided into zones. A zone is defined as a subregion within a field that has similar characteristics in terms of crop management (Klompenburg et al., 2020). Let Z be a set of zones: $Z = \{z_1, z_2, \dots, z_n\}$. For each zone z_i by year, the following features are grouped into two categories: soil data and weather data. The soil data contains information on the soil tests conducted in the agricultural zones. These soil tests are infrequent (typically every 3 to 4 years) because they are expensive and also because the values do not vary substantially over relatively short periods of time. Since we are using a time frame of one year per instance, we assume that the soil properties do not change meaningfully in that time, and so treat them as constants. If a zone does not have a soil test in a given year, the most recent test done in the same zone in the year before is used to map the soil. However, weather data are a collection of a series of time points. An instance I in the data set can be represented as follows:

$$I = [Year, ZoneID, S, W, T, Y] \quad (1)$$

where $S = \{s_1, s_2, \dots, s_x\}$ is the set of x soil variables. $W = \{w_1, w_2, \dots, w_y\}$ is the set of y weather variables, T is time in weeks, and Y is yield.

Figure 1 shows the attributes of the soil and weather data used in this study from 2013 to 2018. It comprises soil nutrients (P, K, Mg), physical properties (soil type, stone content), chemical properties (organic matter, CaCO_3 , pH), and the yield for a zone along with the sowing and harvest dates. The numerous soil type classifications in the data set are

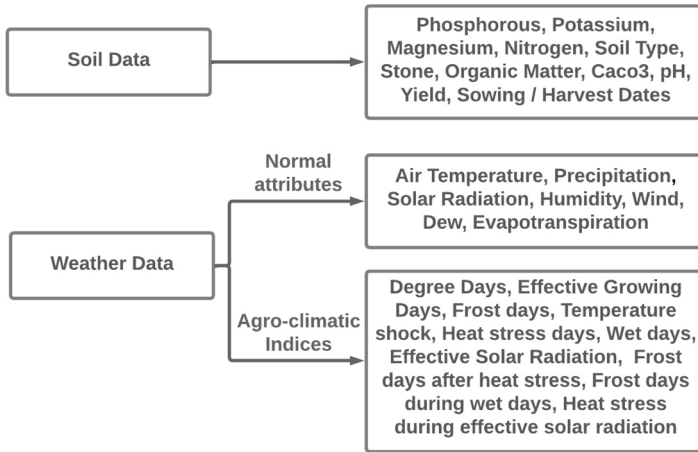


Fig. 1 Data description

shallow, medium, deep clay and deep fertile; stone content is stoneless, low, moderate, high and gravel; organic matter is low, moderate and very high; and CaCO_3 is slightly calc, calc, extremely calc and potentially acidic. Furthermore, weather data include air temperature, precipitation, solar radiation, humidity, wind, dew, and evapotranspiration. More details of the data are presented in this study (Bansal et al., 2022) by the same authors.

4 Data pre-processing

The data pre-processing step is crucial in our study for crop yield prediction, as it involves cleaning and transforming raw data into a format that can be used by machine learning algorithms. It is necessary to remove inconsistencies, missing values, and errors in the data, which could lead to incorrect predictions. Also, it involves transforming the data into a format that is suitable for analysis, such as normalisation or encoding categorical variables.

In this study, after pre-processing soil and weather data sets, those weather features that are calculated by week are integrated with soil features to form a diverse and heterogeneous data set. The features are then represented weekly starting from the 1st week up to the r th week followed by filtering them to focus on the growth period, specifically from the 17th week to the harvest period. This time frame captures the period that has the most significant effect on the yield of the winter wheat crop.

Figure 2 illustrates the pre-processing steps of the soil and weather data. Weather features are calculated both weekly and on the basis of the winter wheat growth stages. The winter wheat growth stages are as follows. Foundation, Construction, Production. All weather features are calculated according to the sowing/harvest dates of a zone/year. These weather features determined weekly are integrated with the soil data after undergoing a principal component analysis. Meanwhile, the weather features derived from winter wheat growth stages pertain to agro-climatic indices that are infrequently observed, leading to an imbalanced data set if calculated by week. When calculating weather features by growth stage, the data can be better balanced, ensuring that the model can learn better.

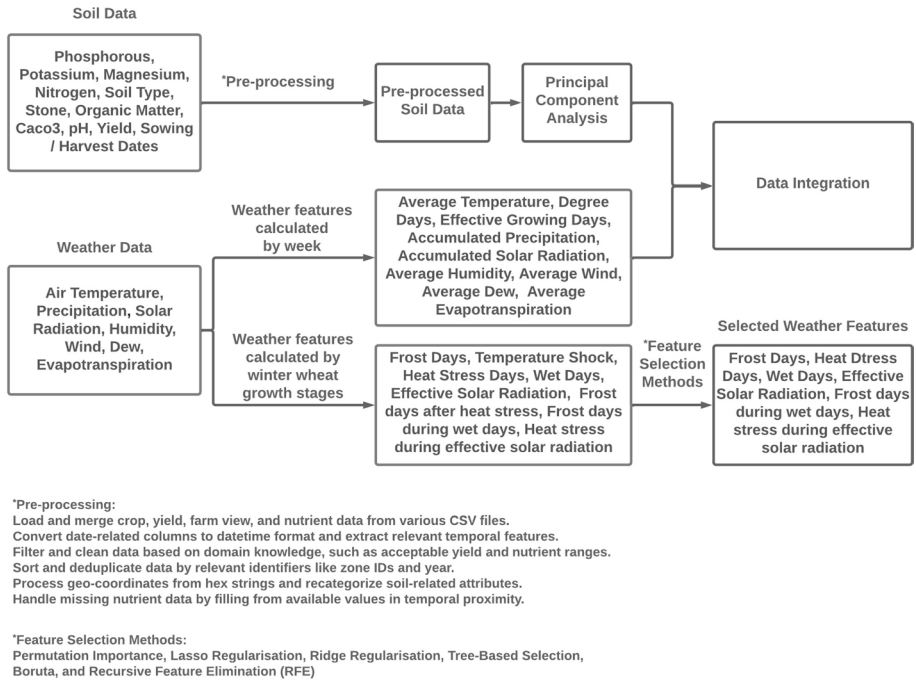


Fig. 2 Data pre-processing

The formulas of all pre-processed weather features are listed below:

- $T_{avg} = \text{mean}((T), \text{week})$ where T_{avg} is the average temperature and T is temperature.
- $DD = \text{sum}(\max(0, (T_{max} + T_{min})/2), \text{week})$ where DD is degree days.
- $EGD = \text{count}((T_{avg} > 5^\circ\text{C}), \text{week})$ where EGD is effective growing days.
- $AP = \text{sum}((P), \text{week})$ where AP is accumulated precipitation.
- $SR_{sum} = \text{sum}((SR), \text{week})$ where SR_{sum} is the total solar radiation.
- $H_{avg} = \text{mean}((H), \text{week})$ where H_{avg} is average humidity.
- $W_{avg} = \text{mean}((\text{Wind}), \text{week})$ where W_{avg} is average wind speed.
- $D_{avg} = \text{mean}((\text{Dew}), \text{week})$ where D_{avg} is average dew point.
- $EV T_{avg} = \text{mean}((EV T), \text{week})$ where $EV T_{avg}$ is average evapotranspiration.
- $FD = \text{count}((T_{min} \leq 0^\circ\text{C}), \text{GrowthStage})$ where FD is the total number of frost days.
- $TS = \text{count}(((T_{avg} > 5^\circ\text{C}) \& (T_{min} > 0^\circ\text{C}) \geq 10) \& (T_{min} < 0^\circ\text{C}), \text{GrowthStage})$ where TS is temperature shock.
- $HSD = \text{count}((T_{max} > 25^\circ\text{C}), \text{GrowthStage})$ where HSD is the total number of heat stress days.
- $WD = \text{count}((P > 10 \text{ mm}), \text{GrowthStage})$ where WD is the total number of wet days.
- $ESR = \text{count}((SR > 5^\circ\text{C}), \text{GrowthStage})$ where ESR is effective solar radiation.
- $FDAHS = \text{count}(((T_{min} \leq 0^\circ\text{C}) \text{ after } (T_{max} > 25^\circ\text{C})), \text{GrowthStage})$ where $FDAHS$ is the number of frost days after heat stress.
- $FDWD = \text{count}((T_{min} < 0^\circ\text{C}) \& (P > 10 \text{ mm}), \text{GrowthStage})$ where $FDWD$ is the number of frost days during wet days.

- $HSESR = \text{count}((T_{\max} > 25^{\circ}\text{C}) \& (T_{\text{avg}} > 5^{\circ}\text{C}), \text{GrowthStage})$ where $HSESR$ is the total number of heat stress days during effective solar radiation.

In this study, we also implemented multiple feature selection techniques (Chandrashekar & Sahin, 2014) to identify the most significant predictors in our data set. These methods included permutation importance, lasso regularisation, ridge regularisation, tree-based selection, Boruta, and recursive feature elimination (RFE). These diverse techniques were chosen to ensure a comprehensive evaluation of feature importance and to minimise the risk of missing relevant predictors. Our analysis revealed that different feature selection methods yielded varying results. However, two methods, Boruta and recursive feature elimination (RFE), consistently identified the same subset of features (i.e., Frost days, heat stress days, wet days, effective solar radiation, frost days during wet days, heat stress during effective solar radiation) as the most important. In particular, these findings aligned well with our existing domain knowledge, further supporting the importance of these features.

5 Proposed approach

Figure 3 outlines a methodology for a neural meta model aimed at predicting the yield of the winter wheat crop. In this study, the data set encompasses the years 2013 to 2018. The training set is derived from the years 2013 to 2017, while the 2018 data serve as the test set. A fivefold cross-validation method is employed on the training data to identify the optimal hyperparameters for each model utilised in this research. Following cross-validation, the final model is trained using the best hyperparameters and subsequently evaluated against the test set.

Initially, we employed a range of classical ML models such as decision tree (DT), support vector (SVR), random forest (RF), extra trees (ET), lightGBM (LGBM), gradient boosting (GB), as well as LSTM and its stabilised version designed for heterogeneous time series datasets. These data sets are primarily made up of weekly calculated weather data features and soil data characteristics. Additionally, weather features determined by the growth stage are used to forecast winter wheat crop yield using weighted regression models such as the weighted decision tree (WDT), the weighted random forest (WRF), and weighted gradient boosting (WGB). The yield predictions sourced from the stabilised LSTM for heterogeneous time series (SLHTS) and the weighted ML models are subsequently fed as input features into a neural meta model. This model is configured in variations to accommodate two input features (SLHTS and WDT), three input features (SLHTS, WDT, WRF), and four input features (SLHTS, WDT, WRF, WGB).

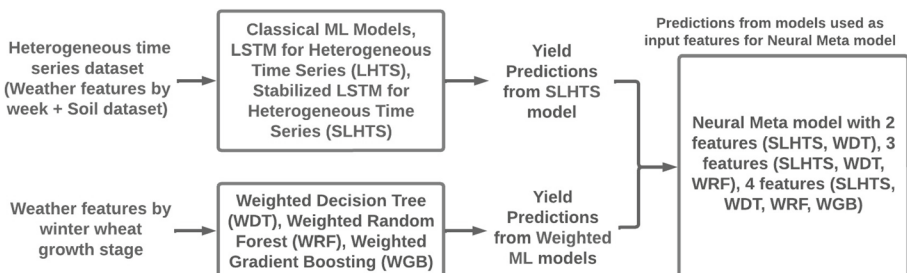


Fig. 3 Proposed approach

WDT, WRF), and four input features (SLHTS, WDT, WRF, WGB). The following subsections provide a comprehensive explanation of each methodology used in our proposed approach.

5.1 LSTM for heterogeneous time series (LHTS)

Figure 4 shows the description of time steps and forward/backward propagation of the LSTM model designed for handling heterogeneous datasets. Each time step represents a week. Since only the winter wheat growth period until harvest weeks are considered in the analysis, the 1st time step starts from the 17th week, the 2nd time step is the 18th week, and the t th time step corresponds to the $16 + t$ th week where t starts from 1. Initially, the cell state and hidden state are initialised to 0 which gets up-dated at each time step based on the input and the previous hidden state.

In the forward time step, weather data for each week of a zone along with its bias is propagated from 17th week to 16th week of that zone. In the backward time step, gradients of weight parameters related to only sequence data, i.e., weather data, are generated. Since nonsequence data, that is, the soil remains constant throughout the year for a zone, it is considered as a bias and its gradients are not computed.

The following are the notations used in Fig. 4:

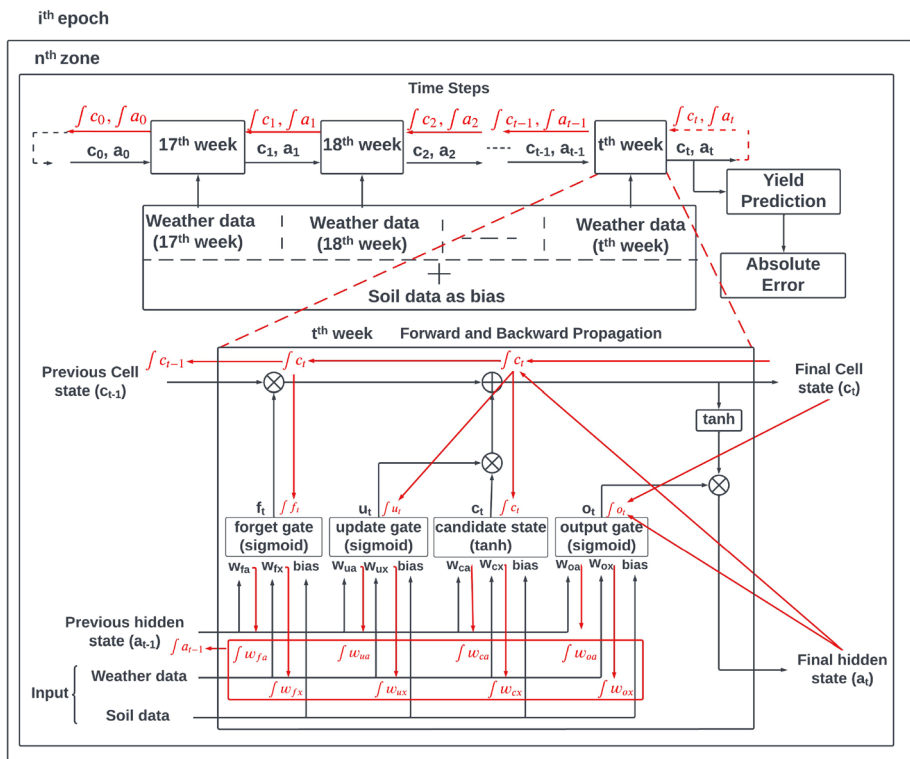


Fig. 4 Architecture of LSTM for heterogeneous time series (LHTS)

f_t , u_t and o_t represents forget, update and output gate, respectively. $\int f_t$, $\int u_t$ and $\int o_t$ are the respective derivatives. w_{fa} , w_{ua} and w_{oa} are the weight parameters related to the hidden state of forget, update and output gate, respectively. $\int w_{fa}$, $\int w_{ua}$, $\int w_{oa}$ are the respective derivatives. w_{fx} , w_{ux} and w_{ox} are the weight parameters related to the weather data from forget, update and output gate, respectively. $\int w_{fx}$, $\int w_{ux}$, and $\int w_{ox}$ are the respective derivatives. a_t , c_t , a_{t-1} , and c_{t-1} represents the previous and current hidden state. Similarly represents the previous and current cell state. $\int a_t$, $\int c_t$, $\int a_{t-1}$, and $\int c_{t-1}$ are the respective derivatives.

In forward propagation, weight parameters of forget gate, update gate, candidate state, output gate, and input comprising of weather and soil data are propagated.

Here, a candidate state updates the cell state on the basis of the input and the previous hidden state. After each time step, the cell state and hidden state are updated. Calculate the gradients of the sequence weather data only by backpropagation over time at time step t using the chain rule. No gradients of soil data are computed during back propagation; instead, it is inputted as a bias. The training and testing steps of the LHTS model are as follows.

- *Data loading, parameter initialisation, and training loop* After reading the datasets and initialising all the model parameters, the model is trained for a specified number of epochs and zones in the training set. For each epoch and zone, the model performs a forward propagation, calculates the absolute error, and updates the model parameters using the Adam optimiser in the backward propagation.
- *Cross-validation scheme* A fivefold cross-validation is used. During each fold, we leave out one year from the training set to serve as the validation set, and the remaining years form the training set.
- *Model training and selection of best hyperparameters* Train the model using different sets of hyperparameters and validate its performance using cross-validation. After the training process is completed, the model will have the best hyperparameters.
- *Final model training and model evaluation* Train the model with the best hyperparameters obtained from cross-validation. The trained model is then evaluated on the test set by predicting the yield and calculating the mean absolute error.

5.2 Error stabilised LSTM for heterogeneous time series (SLHTS)

In this subsection, we dive deeper into the specifics of the SLHTS model, detailing its design, operation, and the distinct characteristics that make it an effective choice for this research. In time-series deep learning models, managing the learning process is of high importance to achieve accurate and reliable predictions. Section 5.1 describes a system architecture for a time series deep learning model suited for heterogeneous data sets and used a conventional backpropagation approach. Although effective, it did not fully take advantage of dynamic control of the learning process based on error stabilisation. Due to this, a novel learning strategy that incorporates the idea of error stabilisation into the training process is presented. The architecture is described in Fig. 5 and applied to heterogeneous data sets. This learning methodology involves monitoring previous n errors during training where n is set to 10 in this study. This parameter can be adjusted according to the characteristics and size of the data used. The error stabilisation condition is invoked when the absolute difference between the current error and the previous error falls within or below the minimum of the previous n errors. Here, current error refers to the absolute

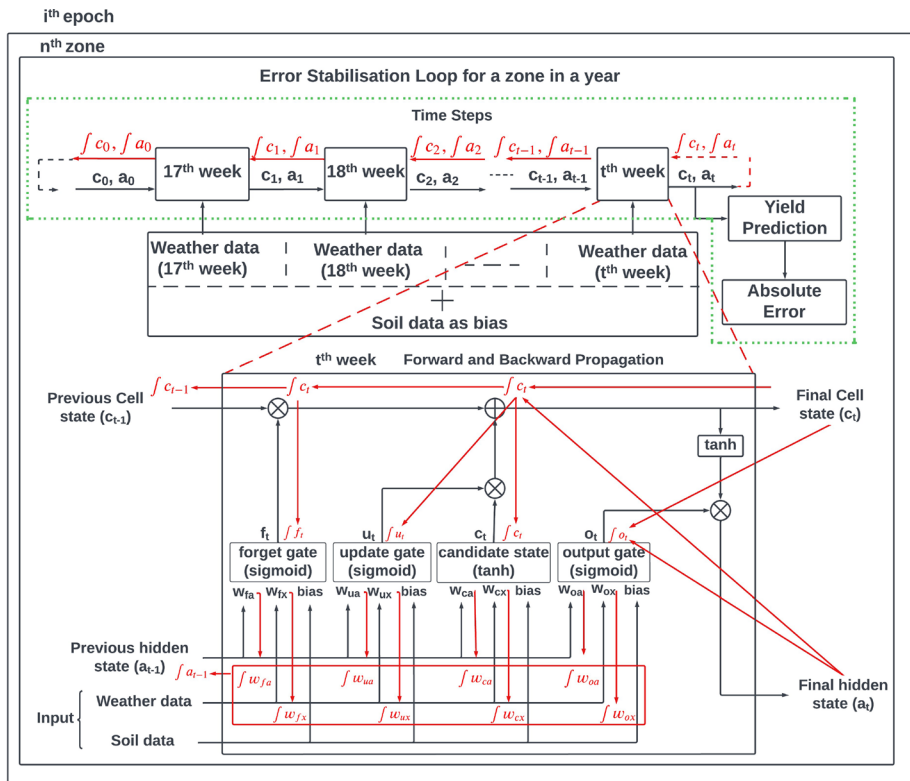


Fig. 5 Architecture of stabilised LSTM for heterogeneous time series (SLHTS)

difference between the actual yield of the zone in a year and the yield predicted by the model for the current iteration of the training phase. The previous error refers to the similar absolute difference, but for the previous iteration in the training phase. This condition states that if the change in error (i.e. the absolute difference between the current and previous errors) does not surpass the smallest change recorded in the recent past (i.e., the minimum of the previous n errors), that means the model learning has stagnated, thereby needing a pause in the training for the current zone in a year. The training and testing steps of the SLHTS model are as follows

- **Data loading, parameter initialisation, and training loop** After reading the datasets and initialising all the model parameters, the model is trained for a specified number of epochs and zones in the training set. For each epoch and zone, the model performs a forward propagation, calculates the absolute error, and updates the model parameters using the Adam optimiser in the backward propagation.
- **Error computation and parameter update** In this methodology, the model performs iterative updates to its parameters, simultaneously calculating a new error with each iteration. This iterative process of parameter update and error computation continues until the difference between the new and old error diminishes to a value smaller than the smallest error observed in the last 10 iterations. Throughout this procedure, the

error and predictions for each zone in a year are stored for subsequent analysis and evaluation.

- *Cross-validation scheme* A fivefold cross-validation is used. During each fold, we leave out one year from the training set to serve as the validation set, and the remaining years form the training set.
- *Model training and selection of best hyperparameters* Train the model using different sets of hyperparameters and validate its performance using cross-validation. After the training process is completed, the model will have the best hyperparameters.
- *Final model training and model evaluation* Train the model with the best hyperparameters obtained from cross-validation. The trained model is then evaluated on the test set by predicting the yield and calculating the mean absolute error.

This approach allows the model to learn until the error stabilises for each zone in a year, providing more fine-tuned control over the learning process.

5.3 Architecture for weighted regression models

This subsection focusses on the model architecture specifically designed for the weather features (i.e., agro-climatic indices) calculated by the winter wheat growth stage. Figure 6 illustrates the architecture of the weighted regression models used in our study. The training and testing steps of this methodology are as follows:

- *Data loading, and weighted regression models* After reading the datasets, three classical ML regression models, i.e., Decision Tree, Random Forest, and Gradient Boosting are customised to be weighted regression models.
- *Hyperparameters and weights introduction* Customising these regression models involves the application of weights along with model-specific hyperparameters to the critical points (lower and higher extreme yield values). Both weights and these critical points are adjustable parameters.
- *Cross-validation scheme* The same fivefold cross-validation is used by leaving out one year from the training set to serve as the validation set and the remaining years to form the training set.
- *Model training and selection of best hyperparameters/weights* Train the model using different sets of hyperparameters and weights. The performance of these models is validated using cross-validation which fine-tunes not only the model hyperparameters but also the weights assigned to the critical data points. After the training process is completed, the regression models will have the best hyperparameters, the best weights, and the critical yield thresholds at which these weights are applied.




Fig. 6 Model architecture for agro-climatic indices dataset




Fig. 7 Neural meta model architecture

- *Final model training and model evaluation* Train the models using the best weights and hyperparameters obtained from cross-validation. The trained model is then evaluated on the test set by predicting the yield and calculating the mean absolute error.

5.4 Neural meta model architecture

This subsection explores the architecture of the neural meta model as stated in Fig. 7. The training and testing steps of a neural meta model are as follows:

- *Data loading, calculation and normalisation of inverse MAE* Obtain train, test predictions, and train MAE from previously trained models (SLHTS and weighted regression models). Then, calculate the inverse MAE and normalise to get weights for each model based on their performance. A model with a lower MAE will have a higher inverse MAE.
- *Weighted predictions and data preparation for a neural meta model* These weights are then multiplied by the model predictions to get weighted predictions. These serve as the input features and are merged to form the data set for a neural meta model.
- *Cross-validation scheme* The same fivefold cross-validation is used by leaving out one year from the training set to serve as the validation set, and the remaining years to form the training set.
- *Model training and selection of best hyperparameters* Train the neural meta model using different sets of hyperparameters, validate its performance using cross-validation. After the training process is completed, the model will have the best hyperparameters.
- *Final model training and model evaluation* Train the model with the best hyperparameters obtained from cross-validation. The trained model is then evaluated on the test set by predicting the yield and calculating the mean absolute error.

6 Experiment setup

- *Computational setup* All experiments were carried out on a machine equipped with Intel Core i7 processor, CPU @ 3.20 GHz 6 cores, and 32 GB RAM. The models were implemented using pandas, numpy, sklearn python libraries and trained on CONTOUR soil dataset and Iteris ClearAg weather dataset described in Sect. 3.
- *Data partitioning* The data set was divided according to years, with training data covering 2013–2017 and test data for 2018. This temporal split ensures that the model is evaluated on unseen future data.
- *Evaluation methodology* Five-fold cross-validation was used in the training set (2013–2017). The performances of the models were evaluated using MAE, RMSE, and R^2

metrics. The primary metrics for this study were MAE and RMSE, with R^2 being secondary. The test set (2018) was used to evaluate the performance of the trained models.

- *SLHTS model* A comprehensive search was conducted to find the best hyperparameters in a validation set using the Adam optimiser. Hyperparameter values included hidden units ranging from 10 to 40, learning rates ranging from 0.001, 0.002, to 0.01, and epochs spanning 10 to 100.
- *Classical and weighted regression ML models* Hyperparameters such as maximum depth (2, 3, 4), minimum sample split (2, 3, 4), and minimum sample leaf (2, 3, 4, 5), learning rate (0.001, 0.01, 0.1) and estimators (100, 200, 300) were considered.
- *Neural meta model* We explored a range of learning rates from 0.001 to 0.01 in increments of 0.001, epochs including 100, 200, 300, 400, and 500, hidden units of 64 and 128, and batch sizes of 16 and 32.
- *Performance evaluation* Consultations with agronomists highlighted that MAE is the most intuitive metric for them due to its expression in metric tons per hectare. Additionally, one-tailed paired t-tests were utilised to assess the statistical significance of performance differences between the various ML models in crop yield predictions.

7 Experimental results

In this section, we dive into the experimental results derived from various ML models deployed to predict winter wheat crop yields. The performance of these models is evaluated based on regression metrics, namely MAE, RMSE, and R^2 . We start with traditional ML approaches, transitioning to an LSTM model designed for heterogeneous datasets that combine soil and weekly calculated weather features. This is succeeded by an LSTM model enhanced for error stabilisation. Subsequently, we present the performance of weighted regression models that use weather features calculated by the winter wheat growth stages based on different harvesting periods in different zones. Concluding the section, results for neural meta model are presented.

Table 1 Classical ML Models

ML Models	Best Hyper-parameters	MAE	RMSE	R2
Decision Tree	$md=2, msl=3, mss=2$	2.57	2.97	0.015
Support Vector	$kernel=rbf, C=1.0, epsilon=0.1$	1.6	1.91	0.03
Random Forest	$md=2, msl=3, mss=3, ne=200$	1.45	1.77	0.002
Extra Trees	$md=2, msl=3, mss=3, ne=200$	1.48	1.79	0.008
Light GBM	$lr=0.01, ne=100$	1.51	1.83	0.008
Gradient Boosting	$md=3, msl=3, mss=2, lr=0.01, ne=100$	1.4	1.69	0.01

The significance of bold is that it shows a better result in comparison with other models


md =max depth, msl =min samples leaf, mss =min samples split, lr =learning rate, ne =n estimators

Table 2 Comparison of SLHTS, LHTS and GB Model. Best Hyper-parameters

Models	Best Hyper-parameters	MAE	RMSE	R2
SLHTS	$lr=0.005, e=10, h=10$	0.98	1.125	0.928
LHTS	$lr=0.005, e=20, h=10$	1.25	1.452	0.591
GB	$md=3, msl=3, mss=2, lr=0.01, ne=100$	1.4	1.69	0.01

The significance of bold is that it shows a better result in comparison with other models

(h =Hidden units, lr =Learning rate, e =Epochs, md =max depth, msl =min samples leaf, mss =min samples split, ne =n estimators)

**Fig. 8** SLHTS versus LHTS versus GB

7.1 Classical ML models

Table 1 shows the performance of traditional ML models on heterogeneous datasets that combine soil and weather features calculated weekly to predict the yield of winter wheat crops. All ML models have very low R values, indicating that they explain a small amount of variance in the dependent variable. Although GB has the best MAE and RMSE, its R^2 is not the highest among models. Given our focus on prediction accuracy, GB is the best performing ML model. In addition, significant differences are observed between SVR and DT ($p \approx 9.61 \times 10^{-37}$), and between RF and SVR ($p \approx 3.53 \times 10^{-8}$). However, comparisons of RF with ET ($p \approx 0.104$), LGBM ($p \approx 0.103$), and GB ($p \approx 0.226$) showed no significant difference in their performance.

7.2 Error stabilised LSTM

Table 2 and Fig. 8 show the performance of Error Stabilised Heterogeneous Time Series (SLHTS), LHTS, GB on heterogeneous datasets that combine soil and weekly calculated weather features. The error stabilised LSTM model (SLHTS) has the best MAE,

RMSE, and R^2 among the models and therefore the best performing model. Additionally, significant differences are observed between SLHTS and LHTS ($p \approx 8.91 \times 10^{-40}$), between SLHTS and GB ($p \approx 1.095 \times 10^{-11}$), and between LHTS and GB ($p \approx 0.0024$).

Table 3 Comparison of Weighted Regression Models

ML Models	Best Hyper-parameters	MAE	RMSE	R2
WDT	$md=2, msl=3, mss=2, lyt=7, wl=1.0, hyt=11, wh=1.2$	1.199	1.142	0.0058
WRF	$md=2, msl=3, mss=3, ne=200, lyt=7, wl=2.6, hyt=11, wh=1.6$	1.23	1.45	0.06
WGB	$md=3, msl=3, mss=2, lr=0.01, ne=100, lyt=8, wl=1.8, hyt=11, wh=2.2$	1.678	1.95	0.039

The significance of bold is that it shows a better result in comparison with other models

WDT=Weighted Decision Tree, WRF=Weighted Random Forest, WGB=Weighted Gradient Boosting, md =max depth, msl =min samples leaf, mss =min samples split, lr =learning rate, ne =nestimators, lyt =low yield threshold, hyt =high yield threshold, wl =weight low, wh =weight high

7.3 Comparison of WDT, WRF, and WGB

Table 3 shows the performance of weighted regression models, Weighted Decision Tree (WDT), Weighted Random Forest (WRF), and Weighted Gradient Boosting (WGB) using weather features calculated by the winter wheat growth stage. While WDT has the best MAE and RMSE, its R^2 is not the highest among the models. It provides the smallest average error and penalizes larger errors the least. Being MAE and RMSE as the primary evaluation metrics in this study, WDT is the best performing model. In addition, significant differences are observed between WDT and WRF ($p \approx 1.98 \times 10^{-5}$), between WDT and WGB ($p \approx 2.14 \times 10^{-14}$), and between WRF and WGB ($p \approx 3.65 \times 10^{-14}$). The output of these weighted regression models fed as input to the neural meta model whose results are presented in the next subsection.

7.4 Neural meta model (NMM) versus error stabilised LSTM for heterogeneous time series (SLHTS)

In Table 4 and Fig. 9, we present the performance of various neural meta models alongside error stabilised LSTM model, SLHTS. The neural meta models, abbreviated as NMM, are differentiated by their input features: NMM 2 with two inputs: predictions

Table 4 Comparison of neural meta model and SLHTS

Models	Best Hyper-parameters	MAE	RMSE	R2
NMM 2	$lr=0.01, e=200, h=32, b=16$	0.824	0.983	0.98
NMM 3	$lr=0.01, e=200, h=32, b=16$	0.854	1.038	0.988
NMM 4	$lr=0.01, e=200, h=32, b=16$	0.869	1.0	0.863
SLHTS	$h=10, lr=0.005, e=10, b=1$	0.98	1.126	0.929

The significance of bold is that it shows a better result in comparison with other models

Best hyper-parameters (lr =Learning rate, e =Epochs, h =Hidden units, b =batch size)

NMM 2, NMM 3, and NMM 4 represent neural meta model with 2, 3, and 4 input features respectively. These features are predictions from weighted regression models and the SLHTS model




Fig. 9 Neural meta model versus SLHTS

from WDT and SLHTS.—NMM 3 with three inputs: predictions from WDT, WRF, and SLHTS.—NMM 4 with four inputs: predictions from WDT, WRF, WGB, and SLHTS. NMM 2 has the best MAE and RMSE, suggesting that it tends to have the smallest average prediction error and also penalizes larger errors the least compared to the other models. NMM 3 has the highest R^2 , implying that it can explain the highest variability in the target variable. Given NMM 2 has the best MAE and RMSE, it is considered as the best performing model. Furthermore, significant differences are observed between NMM 2 and NMM 3 ($p \approx 0.003$), and between NMM 2 and NMM 4 ($p \approx 0.0001$), and between NMM 2 and SLHTS ($p \approx 2.41 \times 10^{-39}$).

8 Ablation study

The importance of this study is shown through an ablation analysis, which provides an understanding of the performance contributions of different modeling techniques for predicting winter wheat crop yield. As evidenced in Table 5, the evaluation emphasises MAE and RMSE as the primary metrics, since they directly reflect the prediction accuracy, which is central to this study. Following are the observations:

Table 5 Comparison of GB, LHTS, SLHTS and NMM models

Models	MAE	RMSE
Gradient BOOSTING	1.4	1.69
LSTM for heterogeneous time series	1.25	1.452
Error stabilised LSTM for heterogeneous time series	0.98	1.125
Neural meta model with 2 input features	0.824	0.983
Neural meta model with 3 input features	0.854	1.038
Neural meta model with 4 input features	0.869	1.0

The significance of bold is that it shows a better result in comparison with other models

- o Traditional ML model gradient boosting serves as a baseline with an MAE of 1.4 and RMSE of 1.69.
- o Moving from LSTM model for heterogeneous time series to error stabilised LSTM results in a noticeable performance improvement, with a drop in MAE from 1.25 to 0.98 and in RMSE from 1.452 to 1.125.
- o Among neural meta models (NMM), which integrate predictions from other models as input features, the model with two input features performs best with the lowest MAE of 0.824 and RMSE of 0.983.

In conclusion, a neural meta model with 2 input features came out to be the best performing model in our analysis, emphasizing the value of meta learning in crop yield prediction.

9 Conclusion

In this study, we worked on using machine learning to predict crop yield, focusing on winter wheat. We developed and tested many ML models, from classical models to advanced neural architectures, trying to find the best way to predict yield. Our experimental tests showed the effectiveness of meta model learning, particularly the neural meta model with 2 input features, which came out as the top performing model. The neural meta model performs well in combining individual model strengths, but it has a static nature, i.e., it does not adjust based on the evolving patterns in the data. A key limitation is that it assigns a fixed weight to the predictions of the individual models based on the performance of each model.

To overcome the identified limitation, our future work will refine the neural meta model to be more adaptable. One strategy involves dynamically adjusting the weights based on the absolute errors of each instance rather than using the overall performance of the individual models. This dynamic adjustment will enable the model to continually adapt to the evolving patterns in the data. In addition, we plan to test the neural meta model with larger datasets to check its robustness and ability to scale.

Authors' contributions The first author (Yogesh Bansal) carried out the implementation of the study, drafted the manuscript, and interpreted the study findings under the supervision and guidance of the co-authors (Dr. David Lillis, Prof. Tahar Kechadi). All authors read and approved the final manuscript.

Funding UCD Foundation,16/SPP/3296,Yogesh Bansal. This research is funded under the SFI Strategic Partnerships Program (16/SPP/3296) and is co-funded by Origin Enterprises Plc.

Availability of data and material Because of the confidentiality agreements with the partnering company and the license under which this study has been conducted, we are restricted from making the data available to the public for the moment.

Code availability Because of the confidentiality agreements with the partnering company and the license under which this study has been conducted, we are restricted from making the code available to the public for the moment.

Declarations

Competing interests There are no competing interests.

Consent to participate Yes, we will participate

Consent for publication The authors give consent for publication of research findings.

References

- Bali, N., & Singla, A. (2021). Deep learning based wheat crop yield prediction model in punjab region of north india. *Applied Artificial Intelligence*, 35(15), 1304–1328.
- Bali, N., & Singla, A. (2022). Emerging trends in machine learning to predict crop yield and study its influential factors: A survey. *Archives of Computational Methods in Engineering*, 29(1), 95–112.
- Bansal, Y., Lillis, D., & Kechadi, M.-T. (2023). A deep learning model for heterogeneous dataset analysis - application to winter wheat crop yield prediction. In J. Abawajy, J. Tavares, L. Kharb, D. Chahal, & A. B. Nassif (Eds.), *Information, communication and computing technology* (pp. 182–194). Springer Nature Switzerland.
- Bansal, Y., Lillis, D., Kechadi, & Tahar, M. (2022). Winter wheat crop yield prediction on multiple heterogeneous datasets using machine learning. In *2022 international conference on computational science and computational intelligence (CSCI)* (pp. 206–212). <https://doi.org/10.1109/CSCI58124.2022.00040>
- Chandrashekar, G., & Sahin, F. (2014). A survey on feature selection methods. *Computers & Electrical Engineering*, 40(1), 16–28.
- Chergui, N., & Kechadi, M. T. (2022). Data analytics for crop management: a big data view. *Journal of Big Data*, 9(1), 1–37.
- Dang, C., Liu, Y., Yue, H., Qian, J. X., & Zhu, R. (2021). Autumn crop yield prediction using data-driven approaches: support vector machines, random forest, and deep neural network methods. *Canadian Journal of Remote Sensing*, 47(2), 162–181.
- Elavarasan, D., Vincent, D. R., Sharma, V., Zomaya, A. Y., & Srinivasan, K. (2018). Forecasting yield by integrating agrarian factors and machine learning models: A survey. *Computers and Electronics in Agriculture*, 155, 257–282.
- Iniyas, S., Varma, V. A., & Naidu, C. T. (2023). Crop yield prediction using machine learning techniques. *Advances in Engineering Software*, 175, 103326.
- Saeed Khaki, Lizhi Wang, and Sotirios V Archontoulis. A cnn-rnn framework for crop yield prediction. *Frontiers in Plant Science*, 10:1750, 2020.
- Mathieu, J. A., & Aires, F. (2018). Assessment of the agro-climatic indices to improve crop yield forecasting. *Agricultural and Forest Meteorology*, 253, 15–30.
- Muruganantham, P., Wibowo, S., Grandhi, S., Samrat, N. H., & Islam, N. (2022). A systematic literature review on crop yield prediction with deep learning and remote sensing. *Remote Sensing*, 14(9), 1990.
- Oikonomidis, A., Catal, C., & Kassahun, A. (2022). Deep learning for crop yield prediction: a systematic literature review. *New Zealand Journal of Crop and Horticultural Science* (pp. 1–26).
- Oikonomidis, A., Catal, C., & Kassahun, A. (2022b). Hybrid deep learning-based models for crop yield prediction. *Applied Artificial Intelligence*, 36(1), 2031822.
- Pravallika, K., Karuna, G., Anuradha, K., Srilakshmi, V. (2021). Deep neural network model for proficient crop yield prediction. In *E3S web of conferences* (Vol. 309, p. 01031. EDP Sciences).
- Sarith Divakar, M., Sudheep Elayidom, M., & Rajesh, R. (2022). Forecasting crop yield with deep learning based ensemble model. *Materials Today: Proceedings*, 58, 256–259.
- Shahhosseini, M., Hu, G., & Archontoulis, S. V. (2020). Forecasting corn yield with machine learning ensembles. *Frontiers in Plant Science*, 11, 1120.
- Shahhosseini, M., Hu, G., Khaki, S., & Archontoulis, S. V. (2021). Corn yield prediction with ensemble cnn-dnn. *Frontiers in Plant Science*, 12, 709008.
- Sharma, S., Rai, S. & Krishnan, N. C. (2020). Wheat crop yield prediction using deep lstm model. arXiv preprint [arXiv:2011.01498](https://arxiv.org/abs/2011.01498).
- Shen, Y., Mercatoris, Benoît, Cao, Z., Kwan, P., Guo, L., Yao, H., & Cheng, Q. (2022). Improving wheat yield prediction accuracy using lstm-rf framework based on uav thermal infrared and multispectral imagery. *Agriculture*, 12(6), 892.

- Srivastava, A. K., Safaei, N., Khaki, S., Lopez, G., Zeng, W., Ewert, F., Gaiser, T., & Rahimi, J. (2022). Winter wheat yield prediction using convolutional neural networks from environmental and phenological data. *Scientific Reports*, 12(1):3215.
- Sun, J., Di, L., Sun, Z., Shen, Y., & Lai, Z. (2019). County-level soybean yield prediction using deep cnn-lstm model. *Sensors*, 19(20), 4363.
- Toomula, S., & Pelluri, S. (2022) An extensive survey of deep learning-based crop yield prediction models for precision agriculture. In *Proceedings of the international conference on cognitive and intelligent computing: ICCIC 2021* (Vol. 1, pp. 1–12. Springer).
- Van Klompenburg, T., Kassahun, A., & Catal, C. (2020). Crop yield prediction using machine learning: A systematic literature review. *Computers and Electronics in Agriculture*, 177, 105709.
- Wang, X., Huang, J., Feng, Q., & Yin, D. (2020). Winter wheat yield prediction at county level and uncertainty analysis in main wheat-producing regions of china with deep learning approaches. *Remote Sensing*, 12(11), 1744.
- Zhang, Y., Liu, J., & Shen, W. (2022b). A review of ensemble learning algorithms used in remote sensing applications. *Applied Sciences*, 12(17), 8654.
- Zhang, Y., Wang, L., Chen, X., Liu, Y., Wang, S., & Wang, L. (2022a). Prediction of winter wheat yield at county level in china using ensemble learning. *Progress in Physical Geography: Earth and Environment*, 46(5), 676–696.

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Springer Nature or its licensor (e.g. a society or other partner) holds exclusive rights to this article under a publishing agreement with the author(s) or other rightsholder(s); author self-archiving of the accepted manuscript version of this article is solely governed by the terms of such publishing agreement and applicable law.

Terms and Conditions

Springer Nature journal content, brought to you courtesy of Springer Nature Customer Service Center GmbH (“Springer Nature”).

Springer Nature supports a reasonable amount of sharing of research papers by authors, subscribers and authorised users (“Users”), for small-scale personal, non-commercial use provided that all copyright, trade and service marks and other proprietary notices are maintained. By accessing, sharing, receiving or otherwise using the Springer Nature journal content you agree to these terms of use (“Terms”). For these purposes, Springer Nature considers academic use (by researchers and students) to be non-commercial.

These Terms are supplementary and will apply in addition to any applicable website terms and conditions, a relevant site licence or a personal subscription. These Terms will prevail over any conflict or ambiguity with regards to the relevant terms, a site licence or a personal subscription (to the extent of the conflict or ambiguity only). For Creative Commons-licensed articles, the terms of the Creative Commons license used will apply.

We collect and use personal data to provide access to the Springer Nature journal content. We may also use these personal data internally within ResearchGate and Springer Nature and as agreed share it, in an anonymised way, for purposes of tracking, analysis and reporting. We will not otherwise disclose your personal data outside the ResearchGate or the Springer Nature group of companies unless we have your permission as detailed in the Privacy Policy.

While Users may use the Springer Nature journal content for small scale, personal non-commercial use, it is important to note that Users may not:

1. use such content for the purpose of providing other users with access on a regular or large scale basis or as a means to circumvent access control;
2. use such content where to do so would be considered a criminal or statutory offence in any jurisdiction, or gives rise to civil liability, or is otherwise unlawful;
3. falsely or misleadingly imply or suggest endorsement, approval, sponsorship, or association unless explicitly agreed to by Springer Nature in writing;
4. use bots or other automated methods to access the content or redirect messages
5. override any security feature or exclusionary protocol; or
6. share the content in order to create substitute for Springer Nature products or services or a systematic database of Springer Nature journal content.

In line with the restriction against commercial use, Springer Nature does not permit the creation of a product or service that creates revenue, royalties, rent or income from our content or its inclusion as part of a paid for service or for other commercial gain. Springer Nature journal content cannot be used for inter-library loans and librarians may not upload Springer Nature journal content on a large scale into their, or any other, institutional repository.

These terms of use are reviewed regularly and may be amended at any time. Springer Nature is not obligated to publish any information or content on this website and may remove it or features or functionality at our sole discretion, at any time with or without notice. Springer Nature may revoke this licence to you at any time and remove access to any copies of the Springer Nature journal content which have been saved.

To the fullest extent permitted by law, Springer Nature makes no warranties, representations or guarantees to Users, either express or implied with respect to the Springer nature journal content and all parties disclaim and waive any implied warranties or warranties imposed by law, including merchantability or fitness for any particular purpose.

Please note that these rights do not automatically extend to content, data or other material published by Springer Nature that may be licensed from third parties.

If you would like to use or distribute our Springer Nature journal content to a wider audience or on a regular basis or in any other manner not expressly permitted by these Terms, please contact Springer Nature at

onlineservice@springernature.com